

WEEK 4

Perform the following DB operations using Cassandra.

1. Create a keyspace by name Employee

```
create keyspace Employee with replication = {  
    ... 'class': 'SimpleStrategy',  
    ... 'replication_factor': 1  
    ... };
```

2. Create a column family by name

Employee-Info with attributes

Emp_Id Primary Key, Emp_Name,
Designation, Date_of_Joining, Salary,
Dept_Name

```
create table employee_info ( Employee_id int PRIMARY KEY, Emp_name text, Designation text,  
Dateofjoining timestamp, Salary double, Dept_name text);
```

3. Insert the values into the table in batch

```
BEGIN BATCH insert into employee_info (Employee_id, Emp_name, Designation,  
Dateofjoining, Salary, Dept_name) values (1, 'Michael Storm', 'HR', '2018-09-22', 200000, 'A')  
insert into employee_info (Employee_id, Emp_name, Designation, Dateofjoining, Salary,  
Dept_name) values (2, 'Stephen Fox', 'Manager', '2019-05-02', 150000, 'B') insert into  
employee_info (Employee_id, Emp_name, Designation, Dateofjoining, Salary, Dept_name)  
values (3, 'David', 'AVP', '2019-03-01', 500000, 'C') insert into employee_info (Employee_id,  
Emp_name, Designation, Dateofjoining, Salary, Dept_name) values (121, 'Ian', 'AVP',  
'2018-03-10', 250000, 'D') APPLY BATCH;
```

4. Update Employee name and Department of Emp-Id 121

```
UPDATE employee_info SET Emp_name = 'Charlie', Dept_name = 'E' WHERE  
Employee_id=121;
```

```
Activities Terminal Apr 16 15:12 bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~  
bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: $ cqlsh  
Connected to Test Cluster at 127.0.0.1:9042  
[cqlsh 0.1.0 | Cassandra 4.1.4 | CQL spec 3.4.6 | Native protocol v5]  
Use HELP for help.  
cqlsh> create keyspace Employee with replication = {  
... 'class': 'SimpleStrategy',  
... 'replication_factor': 1  
... };  
ConfigurationException: Unable to find replication strategy class 'org.apache.cassandra.locator.SimpleStrategy'  
cqlsh> create keyspace Employee with replication = { 'class': 'SimpleStrategy', 'replication_factor': 1 };  
cqlsh> USE Employee;  
cqlsh:employee> create table employee_info(  
... Employee_id PRIMARY KEY,  
...  
cqlsh:employee> create table employee_info ( Employee_id int PRIMARY KEY, Emp_name text, Designation text, Dateofjoining timestamp, Salary double, Dept_name text);  
cqlsh:employee> describe Employee_info  
  
CREATE TABLE employee.employee_info (  
  employee_id int PRIMARY KEY,  
  dateofjoining timestamp,  
  dept_name text,  
  designation text,  
  emp_name text,  
  salary double  
  ) WITH additional_write_policy = '99p'  
  AND bloom_filter_fp_chance = 0.01  
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}  
  AND cdc = false  
  AND comment = ''  
  AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}  
  AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}  
  AND crc_check_chance = 1.0  
  AND default_time_to_live = 0  
  AND extensions = {}  
  AND gc_grace_seconds = 864000  
  AND max_index_interval = 2048  
  AND memtable_flush_period_in_ms = 0  
  AND min_index_interval = 128  
  AND read_repair = 'BLOCKING'  
  AND speculative_retry = '99p';  
cqlsh:employee> describe table employee_info  
  
CREATE TABLE employee.employee_info (  
  employee_id int PRIMARY KEY,  
  dateofjoining timestamp,  
  dept_name text,  
  designation text,  
  emp_name text,  
  salary double  
  ) WITH additional_write_policy = '99p'  
  AND bloom_filter_fp_chance = 0.01  
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}  
  AND cdc = false  
  AND comment = ''  
  AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}  
  AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}  
  AND crc_check_chance = 1.0  
  AND default_time_to_live = 0  
  AND extensions = {}  
  AND gc_grace_seconds = 864000  
  AND max_index_interval = 2048  
  AND memtable_flush_period_in_ms = 0  
  AND min_index_interval = 128  
  AND read_repair = 'BLOCKING'  
  AND speculative_retry = '99p';
```

```
Activities Terminal Apr 16 15:13 bmsccse@bmsccse-HP-Elite-Tower-800-G9-Desktop-PC: ~  
AND extensions = {}  
AND gc_grace_seconds = 864000  
AND max_index_interval = 2048  
AND memtable_flush_period_in_ms = 0  
AND min_index_interval = 128  
AND read_repair = 'BLOCKING'  
AND speculative_retry = '99p';  
cqlsh:employee> insert into employee_info (Employee_id, Emp_name, Designation, Dateofjoining, Salary, Dept_name) values (1, 'Michael Storm', 'HR', '2018-09-22', 200000, 'A')  
... insert into employee_info (Employee_id, Emp_name, Designation, Dateofjoining, Salary, Dept_name) values (2, 'Stephen Fox', 'Manager', '2019-05-02', 150000, 'B')  
... insert into employee_info (Employee_id, Emp_name, Designation, Dateofjoining, Salary, Dept_name) values (3, 'David', 'AVP', '2019-03-01', 500000, 'C')  
... insert into employee_info (Employee_id, Emp_name, Designation, Dateofjoining, Salary, Dept_name) values (121, 'Ian', 'AVP', '2018-03-10', 250000, 'D')  
... APPLY BATCH;  
SyntaxException: line 2:6 mismatched input 'insert' expecting EOF (... '2018-09-22', 200000, 'A')[Insert]...  
cqlsh:employee> BEGIN BATCH insert into employee_info (Employee_id, Emp_name, Designation, Dateofjoining, Salary, Dept_name) values (1, 'Michael Storm', 'HR', '2018-09-22', 200000, 'A') insert into employee_info (Employee_id, Emp_name, Designation, Dateofjoining, Salary, Dept_name) values (2, 'Stephen Fox', 'Manager', '2019-05-02', 150000, 'B') insert into employee_info (Employee_id, Emp_name, Designation, Dateofjoining, Salary, Dept_name) values (3, 'David', 'AVP', '2019-03-01', 500000, 'C') insert into employee_info (Employee_id, Emp_name, Designation, Dateofjoining, Salary, Dept_name) values (121, 'Ian', 'AVP', '2018-03-10', 250000, 'D') APPLY BATCH;  
cqlsh:employee> select * from employee_info  
... select * from employee_info  
cqlsh:employee> select * from employee_info;  
  
employee_id | dateofjoining | dept_name | designation | emp_name | salary  
-----  
1 | 2018-09-21 18:30:00.000000+0000 | A | HR | Michael Storm | 2e+05  
2 | 2019-05-01 18:30:00.000000+0000 | B | Manager | Stephen Fox | 1.5e+05  
121 | 2018-03-09 18:30:00.000000+0000 | D | AVP | Ian | 2.5e+05  
3 | 2019-02-28 18:30:00.000000+0000 | C | AVP | David | 5e+05  
(4 rows)  
cqlsh:employee> update employee_info SET Emp_name = 'Charlie', Dept = 'E' WHERE Employee_id=121;  
InvalidRequest: Error from server: code=2208 [Invalid query] message="Undefined column name dept in table employee.employee_info"  
cqlsh:employee> update employee_info SET Emp_name = 'Charlie' Dept = 'E' WHERE Employee_id=121;  
SyntaxException: line 1:46 no viable alternative at input 'Dept' (...employee_info SET Emp_name = 'Charlie' [Dept]...)  
cqlsh:employee> update employee_info SET Emp_name = 'Charlie' SET Dept = 'E' WHERE Employee_id=121;  
SyntaxException: line 1:46 no viable alternative at input 'SET' (...employee_info SET Emp_name = 'Charlie' [SET]...)  
cqlsh:employee> UPDATE employee_info SET Emp_name = 'Charlie' SET Dept = 'E' WHERE Employee_id=121;  
SyntaxException: line 1:46 no viable alternative at input 'SET' (...employee_info SET Emp_name = 'Charlie' [SET]...)  
cqlsh:employee> UPDATE employee_info SET Emp_name = 'Charlie' Dept = 'E' WHERE Employee_id=121;  
SyntaxException: line 1:47 no viable alternative at input 'Dept' (...employee_info SET Emp_name = 'Charlie' [Dept]...)  
cqlsh:employee> UPDATE employee_info SET Emp_name = 'Charlie' Dept_name = 'E' WHERE Employee_id=121;  
SyntaxException: line 1:46 no viable alternative at input 'Dept_name' (...employee_info SET Emp_name = 'Charlie' [Dept_name]...)  
cqlsh:employee> UPDATE employee_info SET Emp_name = 'Charlie', Dept_name = 'E' WHERE Employee_id=121;  
cqlsh:employee> select * from employee_info;  
  
employee_id | dateofjoining | dept_name | designation | emp_name | salary  
-----  
1 | 2018-09-21 18:30:00.000000+0000 | A | HR | Michael Storm | 2e+05  
2 | 2019-05-01 18:30:00.000000+0000 | B | Manager | Stephen Fox | 1.5e+05  
121 | 2018-03-09 18:30:00.000000+0000 | E | AVP | Charlie | 2.5e+05  
3 | 2019-02-28 18:30:00.000000+0000 | C | AVP | David | 5e+05  
(4 rows)  
cqlsh:employee> select * from employee_info order by salary;  
InvalidRequest: Error from server: code=2208 [Invalid query] message="ORDER BY is only supported when the partition key is restricted by an EQ or an IN."  
cqlsh:employee> select * from employee_info order by salary ASC;  
InvalidRequest: Error from server: code=2208 [Invalid query] message="ORDER BY is only supported when the partition key is restricted by an EQ or an IN."
```